

# Decision Tree | Assignment

**Question 1: What is a Decision Tree, and how does it work in the context of classification?**

**Answer:**

A Decision Tree is a supervised machine learning algorithm used for classification and regression that models decision-making as a tree-like structure. It works by repeatedly splitting the dataset into smaller, more homogeneous groups based on feature values, enabling clear and interpretable predictions.

How a Decision Tree Works in Classification:

- Root Node:  
Represents the entire dataset and initiates the first split based on the most informative feature.
- Decision Nodes:  
Apply conditional rules (e.g., *Is Age > 30?*) to partition data, optimizing class separation using metrics such as Gini Impurity or Information Gain (Entropy).
- Leaf Nodes:  
Store the final class labels, representing the model's prediction outcome.

Operational Flow:

1. The algorithm selects the best feature to split the data based on impurity reduction.
2. Data is recursively divided into branches until a stopping condition is met (pure nodes or max depth).
3. A new data point traverses the tree based on its feature values and lands in a leaf node that defines its class.

**Question 2: Explain the concepts of Gini Impurity and Entropy as impurity measures. How do they impact the splits in a Decision Tree?**

**Answer:**

**Gini Impurity** and **Entropy** are **impurity measures** used in **Decision Tree algorithms** to evaluate how well a dataset is split at each node. Their core objective is to quantify **how mixed the class labels are** and guide the algorithm toward **cleaner, more informative splits**.

---

## 1. Gini Impurity

- Measures the **probability of incorrectly classifying** a randomly chosen data point if labels were assigned randomly based on class distribution.

- Formula:

$$\text{Gini} = 1 - \sum_{i=1}^n p_i^2$$

where  $p_i$  is the probability of class  $i$ .

#### **Key Characteristics:**

- Lower Gini value → **purer node**
  - Gini = 0 → all samples belong to a single class
  - Computationally efficient and widely used (default in CART)
- 

## **2. Entropy**

- Measures the **level of disorder or uncertainty** in the data.
- Formula:

$$\text{Entropy} = - \sum_{i=1}^n p_i \log_2(p_i)$$

#### **Key Characteristics:**

- Entropy = 0 → completely pure node
  - Higher entropy → higher randomness
  - More sensitive to class distribution changes
- 

## **3. Impact on Decision Tree Splits**

- At each node, the algorithm evaluates all possible splits.
- The split that results in the **maximum reduction in impurity** is selected:
  - **Gini Gain** for Gini Impurity
  - **Information Gain** for Entropy

$$\text{Information Gain} = \text{Parent Impurity} - \text{Weighted Child Impurity}$$

**Question 3: What is the difference between Pre-Pruning and Post-Pruning in Decision Trees? Give one practical advantage of using each.**

**Answer:**

**Pre-Pruning** and **Post-Pruning** are optimization strategies used in **Decision Trees** to control model complexity and prevent overfitting. The key difference lies in **when** the tree growth is restricted.

---

### **Pre-Pruning (Early Stopping)**

- Limits tree growth **during the training phase**.
- The algorithm stops splitting a node if certain conditions are met (e.g., maximum depth, minimum samples per split, minimum impurity decrease).

#### **Practical Advantage:**

- **Faster training and lower computational cost**, making it suitable for large datasets and time-sensitive modeling scenarios.
- 

### **Post-Pruning (Late Pruning)**

- Allows the tree to grow fully and then **prunes unnecessary branches after training** based on validation performance.
- Removes branches that do not contribute significantly to predictive accuracy.

#### **Practical Advantage:**

- **Improved generalization performance**, as pruning decisions are based on actual model behavior rather than early assumptions.

## **Question 4: What is Information Gain in Decision Trees, and why is it important for choosing the best split?**

#### **Answer:**

**Information Gain** is a key metric used in **Decision Tree algorithms** to determine the **most effective feature for splitting the data** at each node. It measures how much **uncertainty (impurity)** in the dataset is reduced after performing a split on a particular feature.

#### **Core Concept:**

Information Gain is calculated using **Entropy**, which quantifies randomness in the data.

$$\text{Information Gain} = \text{Entropy}(\text{Parent}) - \sum \text{Entropy}(\text{Children})$$

#### **Why Information Gain Is Important:**

- Identifies the feature that delivers the **maximum reduction in uncertainty**
- Ensures each split results in **more homogeneous (pure) child nodes**
- Helps build **efficient, compact, and accurate** decision trees

- Directly impacts **model performance and interpretability**

#### **Strategic Value:**

By prioritizing splits with the highest Information Gain, decision trees optimize their **decision-making pathway**, reduce overfitting risks, and accelerate convergence toward meaningful outcomes.

In summary, Information Gain acts as a **decision-quality accelerator**, enabling the model to choose splits that maximize clarity, predictive strength, and overall analytical efficiency.

#### **Question 5: What are some common real-world applications of Decision Trees, and what are their main advantages and limitations?**

#### **Answer:**

**Decision Trees** are widely adopted across industries due to their **interpretability, flexibility, and decision-centric structure**. They are used for both **classification and regression** tasks in real-world scenarios.

---

#### **Common Real-World Applications of Decision Trees**

##### **1. Finance & Banking**

- Credit risk assessment
- Loan approval and fraud detection
- Helps institutions make **rule-based, auditable decisions**

##### **2. Healthcare**

- Disease diagnosis and treatment recommendation
- Patient risk stratification
- Enables **transparent clinical decision support**

##### **3. Marketing & Sales**

- Customer segmentation
- Churn prediction and targeted campaigns
- Drives **data-backed personalization strategies**

##### **4. Human Resources**

- Employee performance evaluation
- Attrition prediction
- Supports **talent analytics and workforce planning**

## 5. Manufacturing & Operations

- Fault detection and quality control
  - Process optimization
  - Improves **operational efficiency and reliability**
- 

### Main Advantages of Decision Trees

- **Easy to Understand and Interpret**  
Results are explainable to both technical and non-technical stakeholders.
  - **Minimal Data Preprocessing**  
No strict requirement for data normalization or scaling.
  - **Handles Both Numerical and Categorical Data**  
Highly versatile across diverse datasets.
  - **Captures Non-Linear Relationships**  
Effective in modeling complex decision boundaries.
- 

### Main Limitations of Decision Trees

- **Prone to Overfitting**  
Especially when the tree becomes too deep and complex.
- **Sensitive to Small Data Changes**  
Minor variations can lead to significantly different tree structures.
- **Lower Predictive Accuracy Compared to Ensembles**  
Single trees often underperform compared to Random Forests or Gradient Boosting.
- **Bias Toward Dominant Features**  
Can favor attributes with more levels.

### Dataset Info:

- **Iris Dataset for classification tasks (sklearn.datasets.load\_iris() or provided CSV).**
- **Boston Housing Dataset for regression tasks (sklearn.datasets.load\_boston() or provided CSV).**

### Question 6: Write a Python program to:

- **Load the Iris Dataset**
- **Train a Decision Tree Classifier using the Gini criterion**

- Print the model's accuracy and feature importances (Include your Python code and output in the code box below.)

**Answer:**

```
from sklearn.datasets import load_iris  
from sklearn.model_selection import train_test_split  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.metrics import accuracy_score  
  
iris = load_iris()  
X = iris.data  
y = iris.target  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
  
model = DecisionTreeClassifier(criterion="gini", random_state=42)  
model.fit(X_train, y_train)  
  
y_pred = model.predict(X_test)  
  
accuracy = accuracy_score(y_test, y_pred)  
  
print("Model Accuracy:", accuracy)  
  
print("\nFeature Importances:")  
for feature, importance in zip(iris.feature_names, model.feature_importances_):  
    print(f'{feature}: {importance:.4f}')
```

O/P:

Model Accuracy: 1.0

Feature Importances:

sepal length (cm): 0.0000  
sepal width (cm): 0.0167  
petal length (cm): 0.9061  
petal width (cm): 0.0772

**Question 7: Write a Python program to:**

- **Load the Iris Dataset**
- **Train a Decision Tree Classifier with max\_depth=3 and compare its accuracy to a fully-grown tree.**

**(Include your Python code and output in the code box below.)**

**Answer:**

```
from sklearn.datasets import load_iris  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score
```

```
iris = load_iris()  
X = iris.data  
y = iris.target
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
dt_limited = DecisionTreeClassifier(max_depth=3, random_state=42)  
dt_limited.fit(X_train, y_train)  
y_pred_limited = dt_limited.predict(X_test)  
accuracy_limited = accuracy_score(y_test, y_pred_limited)
```

```
dt_full = DecisionTreeClassifier(random_state=42)
```

```

dt_full.fit(X_train, y_train)
y_pred_full = dt_full.predict(X_test)
accuracy_full = accuracy_score(y_test, y_pred_full)

print("Decision Tree with max_depth=3 Accuracy:", accuracy_limited)
print("Fully-grown Decision Tree Accuracy:", accuracy_full)

```

O/P :

Decision Tree with max\_depth=3 Accuracy: 1.0

Fully-grown Decision Tree Accuracy: 1.0

**Question 8: Write a Python program to:**

- **Load the Boston Housing Dataset**
- **Train a Decision Tree Regressor**
- **Print the Mean Squared Error (MSE) and feature importances (Include your Python code and output in the code box below.)**

**Answer:**

```

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error
import pandas as pd

```

```

data = fetch_california_housing()
X = pd.DataFrame(data.data, columns=data.feature_names)
Y = data.target

```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

```
model = DecisionTreeRegressor(random_state=42)
```

```

model.fit(X_train, Y_train)

Y_pred = model.predict(X_test)

mse = mean_squared_error(Y_test, Y_pred)

print("Mean Squared Error (MSE):", mse)
print("\nFeature Importances:")
for feature, importance in zip(X.columns, model.feature_importances_):
    print(f'{feature}: {importance:.4f}')

```

O/P:

Mean Squared Error (MSE): 0.495235205629094

Feature Importances:

MedInc: 0.5285

HouseAge: 0.0519

AveRooms: 0.0530

AveBedrms: 0.0287

Population: 0.0305

AveOccup: 0.1308

Latitude: 0.0937

Longitude: 0.0829

**Question 9: Write a Python program to:**

- **Load the Iris Dataset**
- **Tune the Decision Tree's max\_depth and min\_samples\_split using GridSearchCV**
- **Print the best parameters and the resulting model accuracy (Include your Python code and output in the code box below.)**

**Answer:**

```
from sklearn.datasets import load_iris
```

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.metrics import accuracy_score

iris = load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

dt = DecisionTreeClassifier(random_state=42)

param_grid = {
    'max_depth': [2, 3, 4, 5, None],
    'min_samples_split': [2, 3, 4, 5]
}

grid_search = GridSearchCV(estimator=dt, param_grid=param_grid, cv=5,
                           scoring='accuracy')
grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

print("Best Parameters:", best_params)
print("Test Set Accuracy:", accuracy)
```

O/P:

Best Parameters: {'max\_depth': 3, 'min\_samples\_split': 2}

Test Set Accuracy: 1.0

**Question 10:** Imagine you're working as a data scientist for a healthcare company that wants to predict whether a patient has a certain disease. You have a large dataset with mixed data types and some missing values. Explain the step-by-step process you would follow to:

- Handle the missing values
- Encode the categorical features • Train a Decision Tree model
- Tune its hyperparameters
- Evaluate its performance And describe what business value this model could provide in the real-world setting.

**Answer:**

### 1. Handle Missing Values

- Identify missing data using methods like isnull() or info().
- Impute missing values depending on the type of data:
  - Numerical features: Replace with mean, median, or use predictive imputation.
  - Categorical features: Replace with mode or create a new category like "Unknown."
- Optionally, drop columns with excessive missingness (>50%) if they are not critical.

```
from sklearn.impute import SimpleImputer
```

```
# Example: numerical imputer
```

```
num_imputer = SimpleImputer(strategy='median')
X_num = num_imputer.fit_transform(X_num)
```

```
# Example: categorical imputer
```

```
cat_imputer = SimpleImputer(strategy='most_frequent')
X_cat = cat_imputer.fit_transform(X_cat)
```

## 2. Encode Categorical Features

1. Convert categorical variables into numeric representations suitable for ML.

2. **One-Hot Encoding:** For nominal categories with no order.

3. **Label Encoding:** For ordinal categories with inherent order.

```
from sklearn.preprocessing import OneHotEncoder
```

```
encoder = OneHotEncoder(sparse=False, handle_unknown='ignore')
```

```
X_cat_encoded = encoder.fit_transform(X_cat)
```

## 3. Train a Decision Tree Model

- Split the data into **training and testing sets**.
- Initialize a **Decision Tree Classifier** and train on the preprocessed data.

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
X_train, X_test, y_train, y_test = train_test_split(X_final, y, test_size=0.2,  
random_state=42)
```

```
model = DecisionTreeClassifier(random_state=42)
```

```
model.fit(X_train, y_train)
```

## 4. Tune Hyperparameters

- Optimize model performance and avoid overfitting using **Grid Search or Randomized Search**.
- Key hyperparameters:
  - `max_depth` → maximum depth of the tree
  - `min_samples_split` → minimum samples required to split a node
  - `min_samples_leaf` → minimum samples at a leaf node
  - `criterion` → “gini” or “entropy” for splitting

```

from sklearn.model_selection import GridSearchCV

param_grid = {
    'max_depth': [3, 5, 7, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

grid_search = GridSearchCV(model, param_grid, cv=5, scoring='accuracy')
grid_search.fit(X_train, y_train)

best_model = grid_search.best_estimator_

```

## 5. Evaluate Performance

- Use metrics suitable for classification:
  - **Accuracy** → overall correctness
  - **Precision & Recall** → especially important for healthcare (false positives vs false negatives)
  - **F1-Score** → balance between precision and recall
  - **ROC-AUC** → probability ranking performance

```

from sklearn.metrics import classification_report, roc_auc_score

y_pred = best_model.predict(X_test)
print(classification_report(y_test, y_pred))

y_prob = best_model.predict_proba(X_test)[:,1]
roc_auc = roc_auc_score(y_test, y_prob)
print("ROC-AUC:", roc_auc)

```

## **6. Business Value in Real-World Healthcare**

- **Early detection:** Helps identify high-risk patients before severe symptoms appear.
- **Resource allocation:** Hospitals can prioritize testing, treatments, or specialist referrals.
- **Decision support:** Assists doctors with evidence-based risk assessment.
- **Cost optimization:** Reduces unnecessary testing and interventions while improving patient outcomes.
- **Scalability:** Can integrate with electronic health records for continuous monitoring and predictive analytics.